

# Classification with Linear Models

Losses for linear classification, logistic regression, multiclass classification

Machine Learning and Data Mining, 2023

Presented by: Abdalaziz Al-Maeeni

Prepared by: Artem Maevskiy

National Research University Higher School of Economics



LAMBD A • HSE

September 30, 2023

Can't we just use linear regression for classification?




# Classification with linear regression

Classification:

$$\hat{f}(x) = \text{sign}[\theta^T x]$$

► For binary classification task, assign:


- $y = +1$  for **positive** class
- $y = -1$  for **negative** class



# Classification with linear regression

Classification:

$$\hat{f}(x) = \text{sign}[\theta^T x]$$




- ▶ For binary classification task, assign:
  - $y = +1$  for **positive** class
  - $y = -1$  for **negative** class
- ▶ Solve linear regression for  $\hat{y} = \theta^T x$  with MSE loss

# Classification with linear regression

Classification:

$$\hat{f}(x) = \text{sign}[\theta^T x]$$




- ▶ For binary classification task, assign:
  - $y = +1$  for **positive** class
  - $y = -1$  for **negative** class
- ▶ Solve linear regression for  $\hat{y} = \theta^T x$  with MSE loss
- ▶ Classify with  $\text{sign}[\hat{y}]$

# Classification with linear regression


Classification:

$$\hat{f}(x) = \text{sign}[\theta^T x]$$




- ▶ For binary classification task, assign:
  - $y = +1$  for **positive** class
  - $y = -1$  for **negative** class
- ▶ Solve linear regression for  $\hat{y} = \theta^T x$  with MSE loss
- ▶ Classify with  $\text{sign}[\hat{y}]$
- ▶ Any problems with this approach?

# Classification with linear regression



- May face problems when classes are unbalanced or have different spread

# Classification with linear regression



MSE loss makes the model **avoid high residuals**

at a price of **pushing the decision boundary** towards the class with higher spread

Can we find a better loss function?



# Classification loss functions




# 0-1 Loss

- Probability of an error

$$\mathcal{L}_{0-1} = \frac{1}{N} \sum_{i=1 \dots N} \mathbb{I}(\theta^T x_i \cdot y_i < 0)$$

$$y_i \in \{-1, +1\}$$




# 0-1 Loss

- Probability of an error

$$\mathcal{L}_{0-1} = \frac{1}{N} \sum_{i=1 \dots N} \mathbb{I}(\theta^T x_i \cdot y_i < 0)$$

$$y_i \in \{-1, +1\}$$




- Can't optimize **piecewise constant** function with gradient-based methods\*

\*other techniques exist (still quite limited)

# Margin

$$M = \theta^T x \cdot y$$


$$\mathcal{L}_{0-1} = \frac{1}{N} \sum_{i=1 \dots N} \mathbb{I}(\theta^T x_i \cdot y_i < 0)$$


margin

$M > 0$  – correct classification

$M < 0$  – incorrect classification

# Upper bounds on 0-1 loss



Instead of optimizing the 0-1 loss we can optimize a **differentiable upper bound**

# Logistic Regression



# Idea

- ▶ Let's model the **class probabilities**

$$P(y = +1 \mid x) = \hat{f}_{\theta}(x)$$
$$P(y = -1 \mid x) = 1 - \hat{f}_{\theta}(x)$$

# Idea

- ▶ Let's model the **class probabilities**

$$P(y = +1 \mid x) = \hat{f}_{\theta}(x)$$
$$P(y = -1 \mid x) = 1 - \hat{f}_{\theta}(x)$$

- ▶ Fit with **maximum (log) likelihood**



# Idea

- ▶ Let's model the **class probabilities**

$$\text{Likelihood} = \prod_{i=1 \dots N} P(y_i \mid x_i)$$

$$P(y = +1 \mid x) = \hat{f}_{\theta}(x)$$
$$P(y = -1 \mid x) = 1 - \hat{f}_{\theta}(x)$$

- ▶ Fit with **maximum (log) likelihood**

# Idea

- ▶ Let's model the **class probabilities**

$$P(y = +1 \mid x) = \hat{f}_\theta(x)$$
$$P(y = -1 \mid x) = 1 - \hat{f}_\theta(x)$$

- ▶ Fit with **maximum (log) likelihood**

$$\begin{aligned}\text{Likelihood} &= \prod_{i=1 \dots N} P(y_i \mid x_i) \\ &= \prod_{i=1 \dots N} \left[ \mathbb{I}[y_i = +1] \cdot \hat{f}_\theta(x_i) \right. \\ &\quad \left. + \mathbb{I}[y_i = -1] \cdot (1 - \hat{f}_\theta(x_i)) \right]\end{aligned}$$

# Idea

- ▶ Let's model the **class probabilities**

$$\begin{aligned}P(y = +1 \mid x) &= \hat{f}_\theta(x) \\P(y = -1 \mid x) &= 1 - \hat{f}_\theta(x)\end{aligned}$$

$$\begin{aligned}\text{Likelihood} &= \prod_{i=1 \dots N} P(y_i \mid x_i) \\&= \prod_{i=1 \dots N} \left[ \mathbb{I}[y_i = +1] \cdot \hat{f}_\theta(x_i) \right. \\&\quad \left. + \mathbb{I}[y_i = -1] \cdot (1 - \hat{f}_\theta(x_i)) \right]\end{aligned}$$

- ▶ Fit with **maximum (log) likelihood**

$$\theta = \operatorname{argmax}_{\theta} \sum_{i=1 \dots N} \left[ \mathbb{I}[y_i = +1] \cdot \log \hat{f}_\theta(x_i) + \mathbb{I}[y_i = -1] \cdot \log(1 - \hat{f}_\theta(x_i)) \right]$$

# Idea

- ▶ Let's model the **class probabilities**

$$\begin{aligned}P(y = +1 \mid x) &= \hat{f}_\theta(x) \\P(y = -1 \mid x) &= 1 - \hat{f}_\theta(x)\end{aligned}$$

$$\begin{aligned}\text{Likelihood} &= \prod_{i=1 \dots N} P(y_i \mid x_i) \\&= \prod_{i=1 \dots N} \left[ \mathbb{I}[y_i = +1] \cdot \hat{f}_\theta(x_i) \right. \\&\quad \left. + \mathbb{I}[y_i = -1] \cdot (1 - \hat{f}_\theta(x_i)) \right]\end{aligned}$$

- ▶ Fit with **maximum (log) likelihood**

$$\theta = \operatorname{argmax}_{\theta} \sum_{i=1 \dots N} \left[ \mathbb{I}[y_i = +1] \cdot \log \hat{f}_\theta(x_i) + \mathbb{I}[y_i = -1] \cdot \log (1 - \hat{f}_\theta(x_i)) \right]$$

- ▶ Predict the class with **highest probability**\*

\*more generally: find a probability threshold suitable for your problem


# Linear probability model

- ▶ How to map the linear model output to a probability value in  $[0, 1]$ ?

# Linear probability model

- ▶ How to map the linear model output to a probability value in  $[0, 1]$ ?
- ▶ Common choice – **sigmoid function**:

$$\sigma(\chi) = \frac{1}{1 + e^{-\chi}}$$




# Linear probability model

- ▶ How to map the linear model output to a probability value in  $[0, 1]$ ?
- ▶ Common choice – **sigmoid function**:

$$\sigma(\chi) = \frac{1}{1 + e^{-\chi}}$$

- ▶ I.e.  $P(y = +1 \mid x) = \sigma(\theta^T x)$




# Linear probability model

- ▶ How to map the linear model output to a probability value in  $[0, 1]$ ?
- ▶ Common choice – **sigmoid function**:

$$\sigma(\chi) = \frac{1}{1 + e^{-\chi}}$$

- ▶ I.e.  $P(y = +1 \mid x) = \sigma(\theta^T x)$
- ▶ Then,  $\theta^T x$  has the meaning of **log odds ratio** between the two classes:

$$\log \frac{P(y = +1 \mid x)}{P(y = -1 \mid x)} = \log \left( \frac{1}{1 + e^{-\theta^T x}} \cdot \frac{1 + e^{-\theta^T x}}{e^{-\theta^T x}} \right) = \theta^T x$$





# Bringing it all together

- Use negative log likelihood as our loss function:

$$\mathcal{L} = - \sum_{i=1 \dots N} \left[ \mathbb{I}[y_i = +1] \cdot \log \hat{f}_{\theta}(x_i) + \mathbb{I}[y_i = -1] \cdot \log(1 - \hat{f}_{\theta}(x_i)) \right]$$

# Bringing it all together

- Use negative log likelihood as our loss function:

$$\mathcal{L} = - \sum_{i=1 \dots N} \left[ \mathbb{I}[y_i = +1] \cdot \log \hat{f}_{\theta}(x_i) + \mathbb{I}[y_i = -1] \cdot \log \left( 1 - \hat{f}_{\theta}(x_i) \right) \right]$$

$$= - \sum_{i=1 \dots N} \left[ \mathbb{I}[y_i = +1] \cdot \log \sigma(\theta^T x_i) + \mathbb{I}[y_i = -1] \cdot \log \sigma(-\theta^T x_i) \right]$$


$$1 - \sigma(x) = \sigma(-x)$$

# Bringing it all together

- Use negative log likelihood as our loss function:

$$\mathcal{L} = - \sum_{i=1 \dots N} \left[ \mathbb{I}[y_i = +1] \cdot \log \hat{f}_{\theta}(x_i) + \mathbb{I}[y_i = -1] \cdot \log \left( 1 - \hat{f}_{\theta}(x_i) \right) \right]$$

$$= - \sum_{i=1 \dots N} \left[ \mathbb{I}[y_i = +1] \cdot \log \sigma(\theta^T x_i) + \mathbb{I}[y_i = -1] \cdot \log \sigma(-\theta^T x_i) \right]$$

$$1 - \sigma(x) = \sigma(-x)$$


$$= - \sum_{i=1 \dots N} \log \sigma(\theta^T x_i \cdot y_i) =$$

# Bringing it all together

- Use negative log likelihood as our loss function:

$$\mathcal{L} = - \sum_{i=1 \dots N} \left[ \mathbb{I}[y_i = +1] \cdot \log \hat{f}_\theta(x_i) + \mathbb{I}[y_i = -1] \cdot \log \left( 1 - \hat{f}_\theta(x_i) \right) \right]$$

$$= - \sum_{i=1 \dots N} \left[ \mathbb{I}[y_i = +1] \cdot \log \sigma(\theta^T x_i) + \mathbb{I}[y_i = -1] \cdot \log \sigma(-\theta^T x_i) \right]$$

$$1 - \sigma(x) = \sigma(-x)$$


$$= - \sum_{i=1 \dots N} \log \sigma(\theta^T x_i \cdot y_i) = \sum_{i=1 \dots N} \log \left( 1 + e^{-\theta^T x_i \cdot y_i} \right)$$

# Bringing it all together

- Use negative log likelihood as our loss function:

$$\mathcal{L} = - \sum_{i=1 \dots N} \left[ \mathbb{I}[y_i = +1] \cdot \log \hat{f}_{\theta}(x_i) + \mathbb{I}[y_i = -1] \cdot \log \left( 1 - \hat{f}_{\theta}(x_i) \right) \right]$$


$$= - \sum_{i=1 \dots N} \left[ \mathbb{I}[y_i = +1] \cdot \log \sigma(\theta^T x_i) + \mathbb{I}[y_i = -1] \cdot \log \sigma(-\theta^T x_i) \right]$$

$$1 - \sigma(x) = \sigma(-x)$$


$$= - \sum_{i=1 \dots N} \log \sigma(\theta^T x_i \cdot y_i) = \sum_{i=1 \dots N} \log \left( 1 + e^{-\theta^T x_i \cdot y_i} \right)$$


- This can be optimized **numerically**

# Example



- Now the boundary is at the right place

# Example




- ▶ Now the boundary is at the right place
- ▶ Note: when classes are linearly separable for any correct decision boundary

$$\theta \longrightarrow C \cdot \theta, \text{ for some } C > 1 \in \mathbb{R}$$

keeps the boundary at the same place, yet improves the loss

# Example



- ▶ Now the boundary is at the right place
- ▶ Note: when classes are linearly separable for any correct decision boundary


$$\theta \longrightarrow C \cdot \theta, \text{ for some } C > 1 \in \mathbb{R}$$

keeps the boundary at the same place, yet improves the loss

- ▶ ideal fit when sigmoid turns into a step function (at infinitely large  $\theta$ )



# Example



- ▶ When classes overlap the loss has a finite minimum
- ▶ Predicted class probability changes smoothly

# Multiclass Logistic Regression

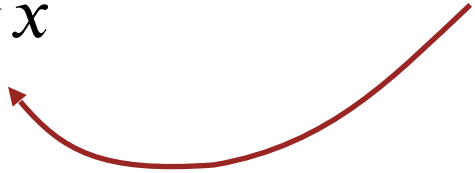


# Multinomial Logistic Regression

- ▶ Similarly to the binary case, we'll model the class probabilities
- ▶ Let's model **unnormalized** class probabilities like this:

$$\tilde{P}(y = k \mid x) = \exp \theta_k^T x$$

Note: now we have  $K$   
parameter vectors

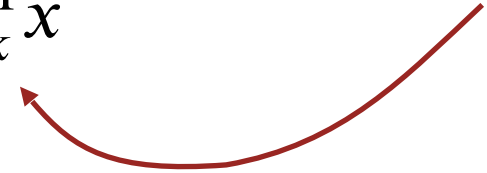


# Multinomial Logistic Regression

- ▶ Similarly to the binary case, we'll model the class probabilities
- ▶ Let's model **unnormalized** class probabilities like this:

$$\tilde{P}(y = k \mid x) = \exp \theta_k^T x$$

Note: now we have  $K$   
parameter vectors



- ▶ Then, the **normalized** probabilities are:

$$P(y = k \mid x) = \frac{\tilde{P}(y = k \mid x)}{\sum_{k'=1 \dots K} \tilde{P}(y = k' \mid x)} = \frac{\exp \theta_k^T x}{\sum_{k'=1 \dots K} \exp \theta_{k'}^T x}$$

- This function is called **softmax** and is commonly used in neural networks

# Multinomial Logistic Regression

- Note that transforming all  $\theta_k \longrightarrow \theta_k + \mathbf{v}$  by some constant vector  $\mathbf{v}$  does not affect the normalized probability

$$\tilde{P}(y = k \mid x) = e^{\theta_k^T x} \longrightarrow e^{\mathbf{v}^T x} \cdot e^{\theta_k^T x} = e^{\mathbf{v}^T x} \cdot \tilde{P}(y = k \mid x)$$

# Multinomial Logistic Regression

- Note that transforming all  $\theta_k \longrightarrow \theta_k + \mathbf{v}$  by some constant vector  $\mathbf{v}$  does not affect the normalized probability

$$\tilde{P}(y = k \mid x) = e^{\theta_k^T x} \longrightarrow e^{\mathbf{v}^T x} \cdot e^{\theta_k^T x} = e^{\mathbf{v}^T x} \cdot \tilde{P}(y = k \mid x)$$

$$P(y = k \mid x) = \frac{\tilde{P}(y = k \mid x)}{\sum_{k'=1 \dots K} \tilde{P}(y = k' \mid x)} \longrightarrow \frac{e^{\mathbf{v}^T x} \cdot \tilde{P}(y = k \mid x)}{\sum_{k'=1 \dots K} e^{\mathbf{v}^T x} \cdot \tilde{P}(y = k' \mid x)} = P(y = k \mid x)$$

# Multinomial Logistic Regression

- Note that transforming all  $\theta_k \longrightarrow \theta_k + \mathbf{v}$  by some constant vector  $\mathbf{v}$  does not affect the normalized probability

$$\tilde{P}(y = k \mid x) = e^{\theta_k^T x} \longrightarrow e^{\mathbf{v}^T x} \cdot e^{\theta_k^T x} = e^{\mathbf{v}^T x} \cdot \tilde{P}(y = k \mid x)$$

$$P(y = k \mid x) = \frac{\tilde{P}(y = k \mid x)}{\sum_{k'=1 \dots K} \tilde{P}(y = k' \mid x)} \longrightarrow \frac{e^{\mathbf{v}^T x} \cdot \tilde{P}(y = k \mid x)}{\sum_{k'=1 \dots K} e^{\mathbf{v}^T x} \cdot \tilde{P}(y = k' \mid x)} = P(y = k \mid x)$$

- This means we can set one of the vectors  $\theta_k$  to 0, e.g. the last one:

$$\theta_K = 0$$

# Multinomial Logistic Regression

- ▶ We now have  $K - 1$  parameter vectors
- ▶ Individual linear outputs  $\theta_k^T x$  now have the meaning of **log odds ratio** between the classes  $k$  and  $K$ :

$$\log \frac{P(y = k | x)}{P(y = K | x)} = \log \frac{\tilde{P}(y = k | x)}{\tilde{P}(y = K | x)} = \log \frac{e^{\theta_k^T x}}{e^0} = \theta_k^T x$$



# Multinomial Logistic Regression

- ▶ Plugging everything into the negative log likelihood we get our loss function:

$$\mathcal{L} = - \sum_{i=1 \dots N} \log \frac{\exp \theta_{y_i}^T x_i}{1 + \sum_{k'=1 \dots K-1} \exp \theta_{k'}^T x_i}$$
$$(\theta_K = 0)$$

- ▶ Again, this can be optimized **numerically**

# Multiclass classification: general approach



# General idea

For a problem with  $K$  classes introduce  $K$  predictors:

$$\hat{f}_k(x): \mathcal{X} \rightarrow \mathbb{R}, \text{ for } k = 1, \dots, K$$

each of which outputs a corresponding **class score**.

Predict the class with the **highest score**:

$$\hat{y}_i = \operatorname{argmax}_k \hat{f}_k(x_i)$$

# Example: binary → multiclass

- ▶ Any binary linear classification model can be converted to multiclass with **one-vs-rest** strategy

^

# Example: binary $\rightarrow$ multiclass

- ▶ Any binary linear classification model can be converted to multiclass with **one-vs-rest** strategy
- ▶ For each class  $k$  train a binary model  $\hat{f}_k(x) = \theta_{(k)}^T x$  separating the given class from all others,  $\hat{y}_{(k)}^{1\text{-vs-rest}} = \text{sign}\left[\hat{f}_k(x)\right]$

# Example: binary $\rightarrow$ multiclass


- ▶ Any binary linear classification model can be converted to multiclass with **one-vs-rest** strategy

- ▶ For each class  $k$  train a binary model  $\hat{f}_k(x) = \theta_{(k)}^T x$  separating the given class from all others,  $\hat{y}_{(k)}^{1-\text{vs}-\text{rest}} = \text{sign}\left[\hat{f}_k(x)\right]$

- ▶ Use the outputs of  $\hat{f}_k$  as class scores for multiclass classification:


$$\hat{y}_i = \underset{k}{\operatorname{argmax}} \hat{f}_k(x_i)$$

# Example



- ▶ Consider the following 3 class problem


# Example



- ▶ “Class-1 VS rest” binary classifier




# Example




- ▶ “Class-2 VS rest” binary classifier

# Example




- ▶ “Class-3 VS rest” binary classifier

# Example




- ▶  $\hat{f}_k(x) = 0$  lines (binary decision boundaries)

# Example



- ▶  $\hat{f}_k(x) = 0$  lines (binary decision boundaries)
- ▶ Adding decision boundaries for  $\hat{y} = \operatorname{argmax}_k \hat{f}_k(x)$


# Example



- ▶ Adding decision boundaries for

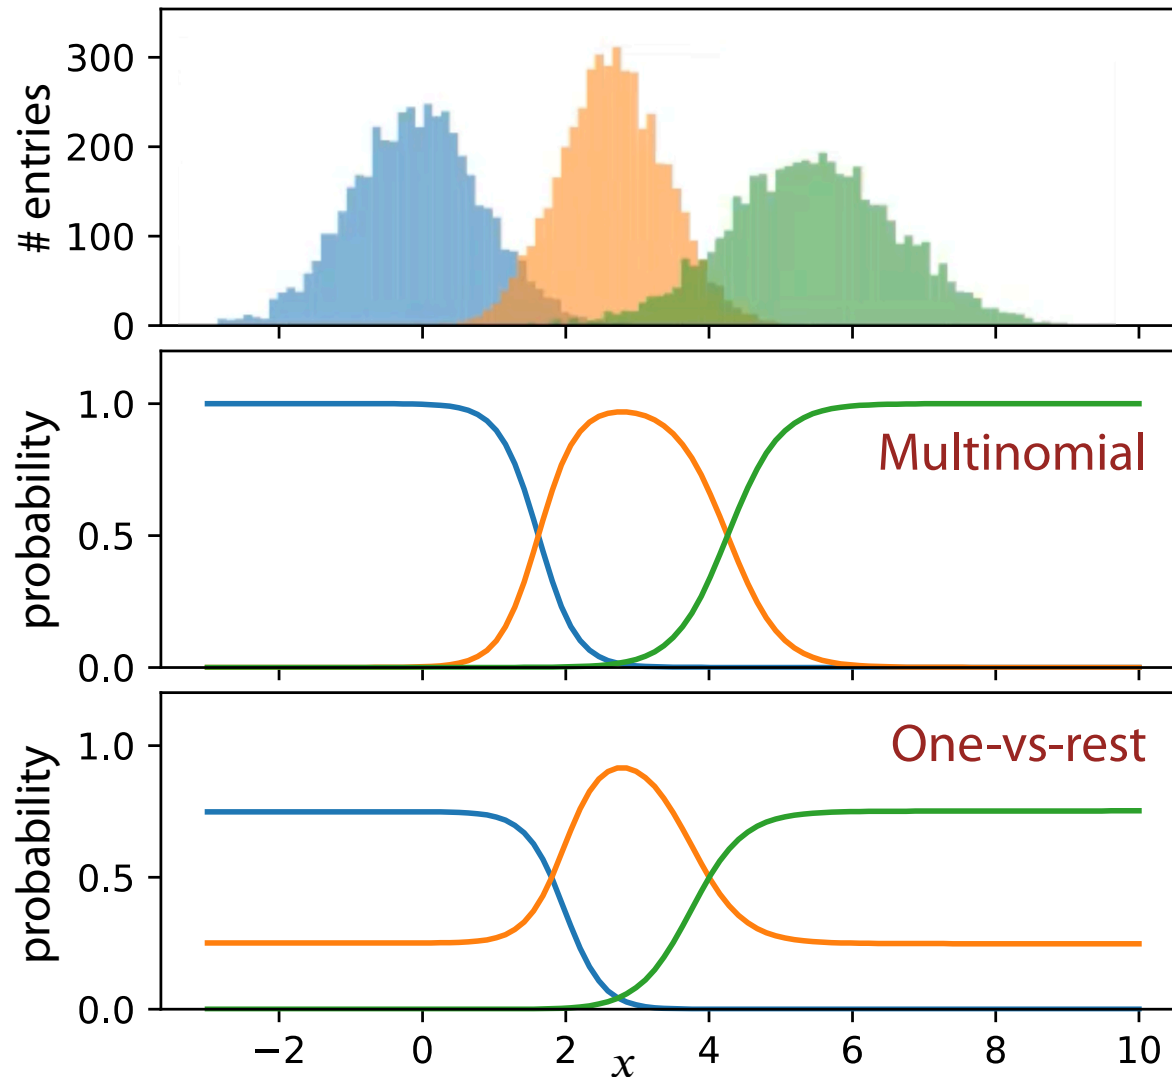
$$\hat{y} = \operatorname{argmax}_k \hat{f}_k(x)$$

# Logistic regression: multinomial or one-vs-rest?



Some of the binary classification tasks not linearly solvable

# Logistic regression: multinomial or one-vs-rest?



Some of the binary classification tasks not linearly solvable

$\Rightarrow$  one-vs-rest results in biased class probabilities

# Summary

- ▶ Classification with linear regression and MSE loss may provide **biased results**



# Summary

- ▶ Classification with linear regression and MSE loss may provide **biased results**
- ▶ 0-1 loss function is better, but is **hard to optimize** directly

# Summary

- ▶ Classification with linear regression and MSE loss may provide **biased results**
- ▶ 0-1 loss function is better, but is **hard to optimize** directly
- ▶ Various **differentiable upper bounds** on 0-1 loss may be used instead

# Summary

- ▶ Classification with linear regression and MSE loss may provide **biased results**
- ▶ 0-1 loss function is better, but is **hard to optimize** directly
- ▶ Various **differentiable upper bounds** on 0-1 loss may be used instead
- ▶ Logistic Regression combines such an upper bound with a **probabilistic model** using the **sigmoid function**

# Summary

- ▶ Classification with linear regression and MSE loss may provide **biased results**
- ▶ 0-1 loss function is better, but is **hard to optimize** directly
- ▶ Various **differentiable upper bounds** on 0-1 loss may be used instead
- ▶ Logistic Regression combines such an upper bound with a **probabilistic model** using the **sigmoid function**
- ▶ Generalizing sigmoid function to a multiclass case yields **softmax function**

# Summary

- ▶ Classification with linear regression and MSE loss may provide **biased results**
- ▶ 0-1 loss function is better, but is **hard to optimize** directly
- ▶ Various **differentiable upper bounds** on 0-1 loss may be used instead
- ▶ Logistic Regression combines such an upper bound with a **probabilistic model** using the **sigmoid function**
- ▶ Generalizing sigmoid function to a multiclass case yields **softmax function**
- ▶ Any binary linear classifier can be adapted to multiclass with the **one-vs-rest strategy**

# Summary

- ▶ Classification with linear regression and MSE loss may provide **biased results**
- ▶ 0-1 loss function is better, but is **hard to optimize** directly
- ▶ Various **differentiable upper bounds** on 0-1 loss may be used instead
- ▶ Logistic Regression combines such an upper bound with a **probabilistic model** using the **sigmoid function**
- ▶ Generalizing sigmoid function to a multiclass case yields **softmax function**
- ▶ Any binary linear classifier can be adapted to multiclass with the **one-vs-rest strategy**
- ▶ Food for thought: how can you mitigate the biased probability problems when using one-vs-rest strategy (as discussed on the previous slide)?

# Thank you!



[al-maeeni@hse.ru](mailto:al-maeeni@hse.ru)



@afdee1c



@AFDEE1C